



UNIVERSITÀ DI PISA

# Computational Intelligence & Deep Learning Project

**A.A. 2022/23**

## Report

Developed by:

- **Domenico Armillotta** - badge : 643020 - email: [d.armillotta@studenti.unipi.it](mailto:d.armillotta@studenti.unipi.it)
- **Stefano Dugo** - badge : 564370 email: [s.dugo@studenti.unipi.it](mailto:s.dugo@studenti.unipi.it)

Link of Google Colab:

[https://colab.research.google.com/drive/1HDUn5\\_bx5Hmeqm3zOZktI\\_gmiK2pzlLx?usp=sharing](https://colab.research.google.com/drive/1HDUn5_bx5Hmeqm3zOZktI_gmiK2pzlLx?usp=sharing)

<b>INTRODUCTION</b>	<b>3</b>
<b>FIRST PART - CNN from Scratch</b>	<b>3</b>
DATASET	3
Prepare the train and test set	4
Handle Imbalanced Dataset	6
Develop the AlexNet Model based Architecture:	6
Evaluate the AlexNet model :	8
Develop the MNIST like Model based architecture:	12
Evaluate the Mnist model :	12
Hyper - Parameter Tuning	17
Evaluate Hyper Mnist Model :	17
<b>SECOND PART - CNN USING VCC19</b>	<b>17</b>
<b>COMPARISON : CNN from Scratch VS CNN with VCC19</b>	<b>23</b>
<b>YOLO</b>	<b>24</b>

## INTRODUCTION

The objective is to create two image classifiers, one developing a CNN from scratch, and the second with a pre-trained keras model but using feature extraction and fine tuning techniques.

Then compare the results.

If there is time left, we will try an implementation using YOL.

## FIRST PART - CNN from Scratch

### DATASET

The German Traffic Sign Benchmark is a multi-class, single-image classification challenge held at the International Joint Conference on Neural Networks (IJCNN) 2011.

Characteristic of dataset :

- More than 40 classes
- More than 50,000 images in total
- Large, lifelike database


Link = [GTSRB - German Traffic Sign Recognition Benchmark | Kaggle](#)

In our project the dataset has been uploaded to google Drive , so that it can be used on google colab.

## Prepare the train and test set


We assigned a textual label to each class.

The plot of the distribution of the data we have is as follows:



The images are organised in folders indicating their class, each one has a different size, but will be resized later.

Here are some examples of street sign images:



We can see that there are images with different gloss, different rotation and so on. Once we have placed the images and labels in a Numpy array, we split the data in the Train folder into Training set and Validation Set according to the 80/20 proportions, while the images in the Test folder we will use to compose the Test set. For labels we use the One Hot Encoding notation.

In the end we will have :

Training = 32k

Validation = 7k

Testing = 18k

### Augmentation:

According to keras documentation when you don't have a large image dataset, it's a good practice to artificially introduce sample diversity by applying random yet realistic transformations to the training images

This helps expose the model to different aspects of the training data while slowing down overfitting.

According to the keras documentation, there are two methods to carry out augmentation, we have chosen to implement the second strategy, i.e. to apply it directly on the train dataset.

With this option, your data augmentation will happen on CPU, asynchronously, and will be buffered before going into the model

The parameters chosen for Augmentation are as follows:

```
rotation_range=10,  
zoom_range=0.15,  
width_shift_range=0.1,  
height_shift_range=0.1,  
shear_range=0.15,  
horizontal_flip=False,  
vertical_flip=False,
```

## Handle Imbalanced Dataset

One way to address class imbalance in a convolutional neural network (CNN) is to use class weights. Class weights are used to adjust the loss function of the model in order to give more weight to the examples in the minority classes, which can help the model to better learn from these classes and improve its overall performance.

In the Keras library, you can specify class weights when compiling a model by using the `class_weight` argument in the `compile()` function. The class weights should be specified as a dictionary that maps class labels to weights.

# Develop the AlexNet Model based Architecture:

## **Architecture:**

This is very similar to the architectures that Alex Krizhevsky advocated in the 2012 for image classification

## **Dropout:**

Dropout technique works by randomly reducing the number of interconnecting neurons within a neural network. At every training step, each neuron has a chance of being left out, or rather, dropped out of the collated contributions from connected neurons.

In the original paper is placed at the end with 0.5 value but [More recent research](#) has shown some value in applying dropout also to convolutional layers, although at much lower levels:  $p=0.1$  or  $0.2$ . Dropout was used after the activation function of each convolutional layer: CONV->RELU->DROP.

## **Convolutional layer:**

A convolution is a mathematical term that describes a dot product multiplication between two sets of elements. Within deep learning the convolution operation acts on the filters/kernels and image data array within the convolutional layer. Therefore a convolutional layer is simply a layer the houses the convolution operation that occurs between the filters and the images passed through a convolutional neural network.

## **Batch Normalisation layer:**

Batch Normalization is a technique that mitigates the effect of unstable gradients within a neural network through the introduction of an additional layer that performs operations on the inputs from the previous layer. The operations standardize and normalize the input values, after that the input values are transformed through scaling and shifting operations.

## **MaxPooling layer:**

Max pooling is a variant of sub-sampling where the maximum pixel value of pixels that fall within the receptive field of a unit within a sub-sampling layer is taken as the output. The max-pooling operation below has a window of  $2 \times 2$  and slides across the input data, outputting an average of the pixels within the receptive field of the kernel.

## **Flatten layer:**

Takes an input shape and flattens the input image data into a one-dimensional array.

## **Dense Layer:**

A dense layer has an embedded number of arbitrary units/neurons within. Each neuron is a perceptron.

A type of activation function that is utilized to derive the probability distribution of a set of numbers within an input vector. The output of a softmax activation function is a vector in which its set of values represents the probability of an occurrence of a class or event. The values within the vector all add up to 1.

[illegible]

Non-trainable params: 960

```
model.compile(loss='categorical_crossentropy',
              optimizer='rmsprop',
              metrics=['accuracy'])
```

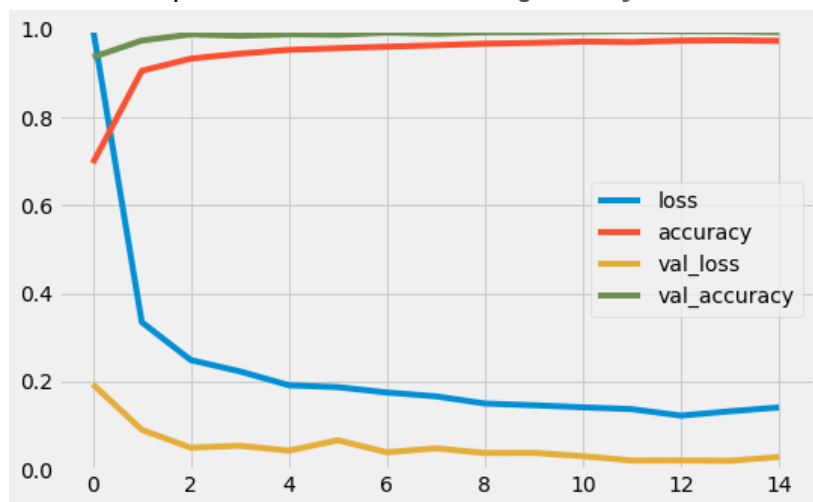
We set also the patience parameters = 10.



```
early_stopping = tf.keras.callbacks.EarlyStopping(
    monitor='val_prc',
    verbose=1,
    patience=10,
    mode='max',
    restore_best_weights=True)
```

## Evaluate the AlexNet model :

Metrics extrapolated from **Model Training History** :



We can see that the model is not in overfitting, because it doesn't have a low error in the training set and a higher error in the testing set.

An attempt was made to combat overfitting by inserting Dropout layers in the model and implementing augmentation of the training set and the early stopping condition.

So we tried to use all the techniques to prevent overfitting in our model.


### Overall accuracy on the Test Set:

395/395 [=====] - 21s 54ms/step

Test Data accuracy: 96.56373713380839

### Confusion Matrix :

Elaborated by submitting the test set




## Report of the model:

For a more detailed report on each class, this graph shows the precision/recall and f1 measure for each model class.

	precision	recall	f1-score	support
0	0.98	1.00	0.99	60
1	0.98	1.00	0.99	720
2	0.99	0.99	0.99	750
3	0.99	0.95	0.97	450
4	1.00	0.97	0.98	660
5	0.95	0.99	0.97	630
6	1.00	0.93	0.97	150
7	0.99	0.99	0.99	450
8	1.00	0.97	0.98	450
9	1.00	0.99	1.00	480
10	1.00	1.00	1.00	660
11	0.96	1.00	0.98	420
12	0.98	0.90	0.94	690
13	0.98	1.00	0.99	720
14	0.86	1.00	0.92	270
15	0.96	1.00	0.98	210
16	0.99	1.00	1.00	150
17	1.00	0.82	0.90	360
18	0.99	0.90	0.94	390
19	1.00	0.53	0.70	60
20	1.00	1.00	1.00	90
21	0.93	0.89	0.91	90
22	1.00	0.79	0.88	120
23	0.80	1.00	0.89	150
24	0.98	0.97	0.97	90
25	0.89	0.99	0.94	480
26	0.98	1.00	0.99	180
27	0.79	0.98	0.87	60
28	1.00	0.96	0.98	150
29	0.74	1.00	0.85	90
30	0.99	0.71	0.83	150
31	1.00	0.99	0.99	270
32	1.00	0.75	0.86	60
33	0.99	1.00	0.99	210
34	0.99	1.00	1.00	120
35	0.99	0.99	0.99	390
36	1.00	0.95	0.97	120
37	1.00	1.00	1.00	60
38	0.91	1.00	0.95	690
39	0.96	0.96	0.96	90
40	0.87	0.98	0.92	90
41	0.98	1.00	0.99	60
42	0.92	0.99	0.95	90
accuracy			0.97	12630
macro avg	0.96	0.95	0.95	12630
weighted avg	0.97	0.97	0.97	12630

### Visual Network Test:


25 random images of the dataset were put to the model test, and we observe how they are classified




Actual=16 || Pred=16



Actual=1 || Pred=1




Actual=38 || Pred=38



Actual=33 || Pred=33




Actual=11 || Pred=11



Actual=38 || Pred=38



Actual=18 || Pred=18



Actual=12 || Pred=12



Actual=25 || Pred=25



Actual=35 || Pred=35



Actual=12 || Pred=12




Actual=7 || Pred=7




Actual=23 || Pred=23



Actual=7 || Pred=7




Actual=4 || Pred=4



Actual=9 || Pred=9




Actual=21 || Pred=21




Actual=20 || Pred=20



Actual=27 || Pred=27



Actual=38 || Pred=38




Actual=4 || Pred=4



Actual=33 || Pred=33



Actual=9 || Pred=9



Actual=3 || Pred=3



Actual=1 || Pred=1


## Develop the MNIST like Model based architecture:

### Architecture:

We will reuse the same general structure of MNIST: our convnet will be a stack of alternated Conv2D (with relu activation) and MaxPooling2D layers.

However, since we are dealing with bigger images and a more complex problem, we will make our network accordingly larger: it will have two more Conv2D + MaxPooling2D stage. This serves both to augment the capacity of the network, and to further reduce the size of the feature maps, so that they aren't overly large when we reach the Flatten layer.

Note that the depth of the feature maps is progressively increasing in the network (from 32 to 256), while the size of the feature maps is decreasing. This is a pattern that you will see in almost all convnets.



Total params: 616,291  
Trainable params: 616,291  
Non-trainable params: 0


### Training Phase:

```
model.compile(loss='categorical_crossentropy',  
              optimizer='rmsprop',  
              metrics=['accuracy'])
```

We have trained the model in 15 epochs, but this is a hyper-parameter that we will tune to in the future

## Evaluate the Mnist model :

Metrics extrapolated from Model Training History :



We can see that the model is not in overfitting , because don't have a low error in the training set and a higher error in the testing set.


An attempt was made to combat overfitting by inserting Dropout layers in the model and implementing augmentation of the training set.

### Overall accuracy on the Test Set:

```
395/395 [=====] - 16s 41ms/step  
Test Data accuracy: 97.79097387173397
```

### Confusion Matrix :

Elaborated by submitting the test set



### Report of the model:


For a more detailed report on each class, this graph shows the precision/recall and f1 measure for each model class.

	precision	recall	f1-score	support
0	0.98	1.00	0.99	60
1	0.97	1.00	0.98	720
2	0.99	0.98	0.98	750
3	0.96	0.96	0.96	450
4	1.00	0.97	0.98	660
5	0.97	0.98	0.98	630
6	0.99	0.98	0.99	150
7	1.00	0.96	0.98	450
8	0.95	0.98	0.97	450
9	1.00	0.99	0.99	480
10	1.00	0.99	0.99	660
11	0.98	0.95	0.97	420
12	1.00	0.98	0.99	690
13	0.99	1.00	1.00	720
14	1.00	1.00	1.00	270
15	0.94	1.00	0.97	210
16	1.00	1.00	1.00	150
17	1.00	0.99	1.00	360
18	0.99	0.96	0.98	390
19	0.71	1.00	0.83	60
20	0.85	1.00	0.92	90
21	0.96	0.72	0.82	90
22	0.98	0.89	0.93	120
23	0.97	1.00	0.98	150
24	0.97	0.93	0.95	90
25	0.97	0.99	0.98	480
26	0.99	1.00	0.99	180
27	0.97	0.93	0.95	60
28	0.96	0.97	0.97	150
29	0.76	1.00	0.86	90
30	0.85	0.72	0.78	150
31	1.00	0.97	0.99	270
32	0.95	1.00	0.98	60
33	0.98	0.99	0.98	210
34	1.00	1.00	1.00	120
35	1.00	0.99	0.99	390
36	0.98	0.98	0.98	120
37	0.97	1.00	0.98	60
38	1.00	1.00	1.00	690
39	0.99	0.96	0.97	90
40	1.00	0.97	0.98	90
41	1.00	1.00	1.00	60
42	0.98	1.00	0.99	90
accuracy			0.98	12630
macro avg	0.96	0.97	0.97	12630
weighted avg	0.98	0.98	0.98	12630

### Visual Network Test:

25 random images of the dataset were put to the model test, and we observe how they are classified







Actual=16 || Pred=16



Actual=1 || Pred=1




Actual=38 || Pred=38



Actual=33 || Pred=33



Actual=11 || Pred=11



Actual=38 || Pred=38



Actual=18 || Pred=18



Actual=12 || Pred=12



Actual=25 || Pred=25



Actual=35 || Pred=35



Actual=12 || Pred=12




Actual=7 || Pred=7




Actual=23 || Pred=23



Actual=7 || Pred=7




Actual=4 || Pred=4



Actual=9 || Pred=9




Actual=21 || Pred=21




Actual=20 || Pred=20



Actual=27 || Pred=27



Actual=38 || Pred=38




Actual=4 || Pred=4



Actual=33 || Pred=33



Actual=9 || Pred=9



Actual=3 || Pred=3



Actual=1 || Pred=1

## Hyper - Parameter Tuning

We tried to fine tuning the hyperparameters in order to improve performance using the Keras library called “keras Tuner”.

This process was only carried out for Mnist, as it was faster to train, alternatively on AlexNet it would have been more or less the same steps.

The optimized parameters were:

the number of training epochs, the learning rate, and the number of units in the dense layers of the network.

All other parameters remained unchanged

This are the result of the search:

```
Best epoch: 37
```

```
The optimal number of units in the first densely-connected layer is 288
```

```
the optimal learning rate for the optimizer is 0.001
```

## Evaluate Hyper Mnist Model :

As can be seen from the graphs, performance has improved slightly, which is due to fine-tuning.

So you can see that it is a very useful step, especially in situations with very large networks









Here are the results:

## Related works :

The results found on the internet regarding the GTSRB dataset, showed similar performance to ours.

On kaggle many of the papers analyzed are superficial on some aspect such as balancing training, or comparing multiple classifiers.

While on github more complete works are present:

1	<b>CNN with 3 Spatial Transformers</b>	99.71%	×	Deep neural network for traffic sign recognition systems: An analysis of spatial transformers and stochastic optimisation methods			2018
2	<b>Sill-Net</b>	99.68%	×	Sill-Net: Feature Augmentation with Separated Illumination Representation			2021
3	<b>MicronNet</b> (fp16)	98.9%	×	MicronNet: A Highly Compact Deep Convolutional Neural Network Architecture for Real-time Embedded Traffic Sign Classification			2018
4	<b>SEER</b> (RegNet10B)	90.71%	✓	Vision Models Are More Robust And Fair When Pretrained On Uncurated Images Without Supervision			2022

## SECOND PART - CNN USING VCC19

### Approach

We used the VGG19 architecture to test a pre-trained network on our dataset through transfer learning. VGG19 is a convolutional neural network (CNN) trained on the ImageNet dataset for image classification. It consists of 19 layers and has been trained to recognize 1000 different classes of objects. For leveraging the network we exploited both feature extraction and fine-tuning.

For the first approach we instantiated the VGG19 convolutional base and then we applied a densely connected classifier on top of it.

We first exploited fast feature extraction without data augmentation: we applied the preprocessing of VGG19 to the input data and then predicted with the convolutional base, then used the densely connected classifier to train, validate and test the model.

Model: "vgg19"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 50, 50, 3)]	0
block1_conv1 (Conv2D)	(None, 50, 50, 64)	1792
block1_conv2 (Conv2D)	(None, 50, 50, 64)	36928
block1_pool (MaxPooling2D)	(None, 25, 25, 64)	0
block2_conv1 (Conv2D)	(None, 25, 25, 128)	73856
block2_conv2 (Conv2D)	(None, 25, 25, 128)	147584
block2_pool (MaxPooling2D)	(None, 12, 12, 128)	0
block3_conv1 (Conv2D)	(None, 12, 12, 256)	295168
block3_conv2 (Conv2D)	(None, 12, 12, 256)	590880
block3_conv3 (Conv2D)	(None, 12, 12, 256)	590880
block3_conv4 (Conv2D)	(None, 12, 12, 256)	590880
block3_pool (MaxPooling2D)	(None, 6, 6, 256)	0
block4_conv1 (Conv2D)	(None, 6, 6, 512)	1180160
block4_conv2 (Conv2D)	(None, 6, 6, 512)	2359808
block4_conv3 (Conv2D)	(None, 6, 6, 512)	2359808
block4_conv4 (Conv2D)	(None, 6, 6, 512)	2359808
block4_pool (MaxPooling2D)	(None, 3, 3, 512)	0
block5_conv1 (Conv2D)	(None, 3, 3, 512)	2359808
block5_conv2 (Conv2D)	(None, 3, 3, 512)	2359808
block5_conv3 (Conv2D)	(None, 3, 3, 512)	2359808
block5_conv4 (Conv2D)	(None, 3, 3, 512)	2359808
block5_pool (MaxPooling2D)	(None, 1, 1, 512)	0
Total params: 20,024,384		
Trainable params: 20,024,384		
Non-trainable params: 0		

We also use feature extraction with a data augmentation layer, freezing all the layers before doing that, and applying the VGG19 convolutional base directly to our model (also applying input scale with *vgg19.preprocess\_input*).

```
#freezing all the layers
conv_base.trainable = False
#Setting trainable to False empties the list of trainable weights of the layer or model
print("Trainable weights after freezing", len(conv_base.trainable_weights))

conv_base.summary()
```

Trainable weights after freezing 0  
Model: "vgg19"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 50, 50, 3)]	0
block1_conv1 (Conv2D)	(None, 50, 50, 64)	1792
block1_conv2 (Conv2D)	(None, 50, 50, 64)	36928
block1_pool (MaxPooling2D)	(None, 25, 25, 64)	0
block2_conv1 (Conv2D)	(None, 25, 25, 128)	73856
block2_conv2 (Conv2D)	(None, 25, 25, 128)	147584
block2_pool (MaxPooling2D)	(None, 12, 12, 128)	0
block3_conv1 (Conv2D)	(None, 12, 12, 256)	295168
block3_conv2 (Conv2D)	(None, 12, 12, 256)	590880
block3_conv3 (Conv2D)	(None, 12, 12, 256)	590880
block3_conv4 (Conv2D)	(None, 12, 12, 256)	590880
block3_pool (MaxPooling2D)	(None, 6, 6, 256)	0
block4_conv1 (Conv2D)	(None, 6, 6, 512)	1180160
block4_conv2 (Conv2D)	(None, 6, 6, 512)	2359808
block4_conv3 (Conv2D)	(None, 6, 6, 512)	2359808
block4_conv4 (Conv2D)	(None, 6, 6, 512)	2359808
block4_pool (MaxPooling2D)	(None, 3, 3, 512)	0
block5_conv1 (Conv2D)	(None, 3, 3, 512)	2359808
block5_conv2 (Conv2D)	(None, 3, 3, 512)	2359808
block5_conv3 (Conv2D)	(None, 3, 3, 512)	2359808
block5_conv4 (Conv2D)	(None, 3, 3, 512)	2359808
block5_pool (MaxPooling2D)	(None, 1, 1, 512)	0
=====		
Total params: 20,024,384		
Trainable params: 0		
Non-trainable params: 20,024,384		

For exploiting fine-tuning, we instantiated the vgg19 convolutional base and unfroze the last 5 layers, then applied the densely connected classifier on top of the network (same model as the feature extraction with data augmentation case). It's not convenient to fine tune all the layers; earlier layers are more generic, while

higher layers are more specialized; also if we fine tune too much layers we risk overfitting.

```
[ ] # Freezing all layers until the fifth from the last
conv_base.trainable = True
for layer in conv_base.layers[:-5]:
    layer.trainable = False
```

```
[ ] conv_base.summary()
```

Model: "vgg19"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 50, 50, 3)]	0
block1_conv1 (Conv2D)	(None, 50, 50, 64)	1792
block1_conv2 (Conv2D)	(None, 50, 50, 64)	36928
block1_pool (MaxPooling2D)	(None, 25, 25, 64)	0
block2_conv1 (Conv2D)	(None, 25, 25, 128)	73856
block2_conv2 (Conv2D)	(None, 25, 25, 128)	147584
block2_pool (MaxPooling2D)	(None, 12, 12, 128)	0
block3_conv1 (Conv2D)	(None, 12, 12, 256)	295168
block3_conv2 (Conv2D)	(None, 12, 12, 256)	590880
block3_conv3 (Conv2D)	(None, 12, 12, 256)	590880
block3_conv4 (Conv2D)	(None, 12, 12, 256)	590880
block3_pool (MaxPooling2D)	(None, 6, 6, 256)	0
block4_conv1 (Conv2D)	(None, 6, 6, 512)	1180160
block4_conv2 (Conv2D)	(None, 6, 6, 512)	2359808
block4_conv3 (Conv2D)	(None, 6, 6, 512)	2359808
block4_conv4 (Conv2D)	(None, 6, 6, 512)	2359808
block4_pool (MaxPooling2D)	(None, 3, 3, 512)	0
block5_conv1 (Conv2D)	(None, 3, 3, 512)	2359808
block5_conv2 (Conv2D)	(None, 3, 3, 512)	2359808
block5_conv3 (Conv2D)	(None, 3, 3, 512)	2359808
block5_conv4 (Conv2D)	(None, 3, 3, 512)	2359808
block5_pool (MaxPooling2D)	(None, 1, 1, 512)	0
=====		
Total params: 20,024,384		
Trainable params: 9,439,232		
Non-trainable params: 10,585,152		

## Parameters

Here we highlight some of the parameters used for the network in the different cases.

### *Loss function*

The loss function we decided to use is the categorical cross entropy.

### *Optimizers*

As optimizer we tried *rmsprop* and the *adam* optimizer:

- **RMSProp:** RMSprop is an optimization algorithm that can be used to train neural networks. It is a variant of stochastic gradient descent that uses moving averages of the squared gradients to scale the learning rates of each parameter. The idea behind RMSprop is to divide the learning rate for a weight by a running average of the magnitudes of the recent gradients for that weight, which can help to prevent the learning rate from becoming too large and causing the optimization process to diverge.
- **Adam:** The Adam optimizer is a popular choice for training neural networks because it is computationally efficient, has good convergence properties, and can be used with little hyperparameter tuning. Adam is a variation of stochastic gradient descent. It uses moving averages of the parameters to provide an running estimate of the second raw moments of the gradients; the term adaptive in the name refers to the fact that the algorithm "adapts" the learning rates of each parameter based on the historical gradient information. This can help the optimization process converge more quickly and with better stability compared to other optimization algorithms.

We also used the *ReduceLROnPlateau* callback to reduce the learning rate of the model since we noticed the validation loss didn't not improve for a number of epochs. This can be useful since we needed to prevent overfitting, as a constantly decreasing learning rate can help the model to converge on a solution; also this can help the model to find a good balance between underfitting and overfitting, as it allows the model to continue learning at a slower pace when the validation loss plateaus.

Also here we exploited early stopping; in some of our tests we noticed that the training stopped early to prevent the validation accuracy to drop.

### *Dropout layer*

We used a Dropout layer for regularization in order to reduce overfitting; we tried different values for this layer and at the end we chose a small value because the size of the dataset was already large enough for our network so if we increased it the performances would decrease; in fact we observed that with larger value for Dropout layers we had a lower training accuracy but also a lower validation accuracy, so it

didn't help preventing overfitting. The value we adopted after some experiments is 0.2.

### *Dense Layer*


In the densely connected classifier built on top of the VGG19 convolutional base, we used the ReLU activation function in the hidden layers. The ReLU (Rectified Linear Unit) activation function is defined as:  $f(x) = \max(0, x)$ . This means that the output of the ReLU function is the input value  $x$  if  $x$  is positive, and 0 if  $x$  is negative. The ReLU activation function has the advantage of being computationally efficient, as well as having a non-saturating gradient, which can facilitate the training process. In the hidden layers of the densely connected classifier, the ReLU activation function can help introduce non-linearity into the model, which can improve its ability to learn complex patterns in the data. Therefore before the output layer, which has a softmax activation function, we decided to use fully connected (dense) layers with the ReLU activation function. The number of input units in this layer was chosen experimentally, starting with low values and at the end we decided to use 1024 units.

## Performances

The accuracy on training and validation sets with fast feature extraction were the following:

### *RMSprop optimizer*

```
[19] #evaluation with history
pd.DataFrame(history.history).plot(figsize=(8, 5))
plt.grid(True)
plt.gca().set_ylim(0, 1)
plt.show()
```



- 4s 4ms/step - loss: 0.0357 - accuracy: 0.9799 - val\_loss: 0.3806 - val\_accuracy: 0.9152 - lr: 0.0010




(Insert here the fast feature extraction test performances + auc)

### *Adam optimizer*

- 3s 4ms/step - loss: 0.0864 - accuracy: 0.9669 - val\_loss: 0.3070 - val\_accuracy: 0.9031 - lr: 1.0000e-04

```
[24] #evaluation with history
pd.DataFrame(history.history).plot(figsize=(8, 5))
plt.grid(True)
plt.gca().set_ylim(0, 1)
plt.show()
```



The accuracy on training and validation sets with fast feature extraction were the following:

loss: 0.0726 - accuracy: 0.9720 - val\_loss: 0.1457 - val\_accuracy: 0.9625 - lr: 0.0010

(Insert here the feature extraction with data augmentation test performances + auc)

(Insert performances for the case for adam optimizer)

The accuracy on training and validation sets with fine tuning were the following:

loss: 0.0763 - accuracy: 0.9628 - val\_loss: 0.1514 - val\_accuracy: 0.9454

(Insert here the fine tuning training and test performances + auc)

(Insert performances for the case for adam optimizer)

(Insert here a little discussion over the performances)

## COMPARISON : CNN from Scratch VS CNN with VCC19

### YOLO

We tried to implement yolo, but unfortunately due to time constraints we could not do it properly.

In fact, we encountered two main problems in the training phase:

1. The labels had to have an edge, and this is usually done with external programmes such as CVAT, Labellmg, VoTT in manual mode, so it would have taken a long time to create a sufficient number of images for our training, validation and test sets.

2. It is generally better to have images of the entire environment rather than just the object itself, especially if you want the model to be able to classify the object in a realistic setting. This is because the context and background of the image can provide valuable information for the model to make an accurate classification.

For example, if you are trying to classify cars in images, it would be helpful for the model to see the car in the context of a street or a parking lot, rather than just a close-up of the car itself. This is because the model can use the surrounding context to understand the shape, size, and appearance of the car, as well as its relationship to other objects in the scene.

So our database was no good, and it was necessary to take a dataset manually or to screen google maps.

