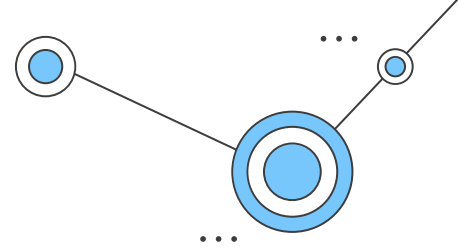


Recognition of traffic signs

Deep Learning project

Domenico Armillotta
Stefano Dugo

Project goal:




IDEA:

The project aims to create an image **classifier** that recognizes traffic signs , given an image. This type of system is particularly useful in applications of self-driving cars or driver assistance in driving.

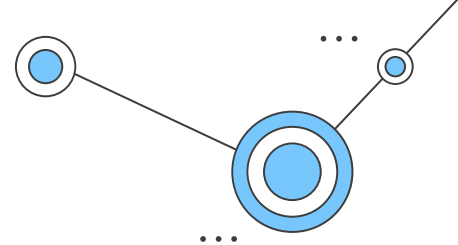
In fact, this specific application is part of ADAS systems that are more generally able to detect, for example, speed limits, but also access and overtaking bans.

IMPLEMENTATION:

The objective is to create **two** image classifiers, one developing a CNN from scratch, and the second with a pre-trained keras model but using feature extraction and fine tuning techniques. Then compare the results.











Related Works:




The results found on the internet regarding the GTSRB dataset, showed similar performance to ours.

On kaggle many of the papers analyzed are superficial on some aspects such as balancing training, or comparing multiple classifiers.

While on github more complete works are present:

1	CNN with 3 Spatial Transformers	99.71%	×	Deep neural network for traffic sign recognition systems: An analysis of spatial transformers and stochastic optimisation methods			2018
2	Sill-Net	99.68%	×	Sill-Net: Feature Augmentation with Separated Illumination Representation			2021
3	MicronNet (fp16)	98.9%	×	MicronNet: A Highly Compact Deep Convolutional Neural Network Architecture for Real-time Embedded Traffic Sign Classification			2018
4	SEER (RegNet10B)	90.71%	✓	Vision Models Are More Robust And Fair When Pretrained On Uncurated Images Without Supervision			2022





01

CNN FROM SCRATCH



1. Dataset


Source: [GTSRB - German Traffic Sign Recognition Benchmark | Kaggle](#)

Class = 43

Dimension : 50k image

Multi-class, single-image classification challenge held at the International Joint Conference on Neural Networks (IJCNN) 2011.

Here are some examples of street sign images:



2. Dataset Splitting & Preparation

Split Phase:

Training = 32k

Validation = 7k

Testing = 18k

Training set and Validation Set according to the 80/20 proportions, while the images in the Test folder we will use to compose the Test set.

Augmentation

it's a good practice to artificially introduce sample diversity by applying random yet realistic transformations. According to the keras documentation, there are two methods to carry out augmentation, we have chosen to implement the second strategy, i.e. to apply it directly on the train dataset.



Handle Imbalanced Dataset

One way to address class imbalance in a convolutional neural network (CNN) is to use class weights. Class weights are used to adjust the loss function of the model in order to give more weight to the examples in the minority classes

In Keras library was used the `class_weight` argument in the `compile()` function.



AlexNet Model Based

Architecture:

- **Dropout** : randomly reducing the number of interconnecting neurons
- **Convolutional layer**: houses the convolution operation that occurs between the filters and the images passed
- **Batch Normalisation layer**: technique that mitigates the effect of unstable gradients
- **MaxPooling layer**: max pixel value of pixels that fall in the receptive field is taken as the output
- **Flatten layer**: flattens the input image data into a one-dimensional array.
- **Dense Layer**: A dense layer has an embedded number of arbitrary units/neurons
- **Softmax Activation Function**: a type of activation function



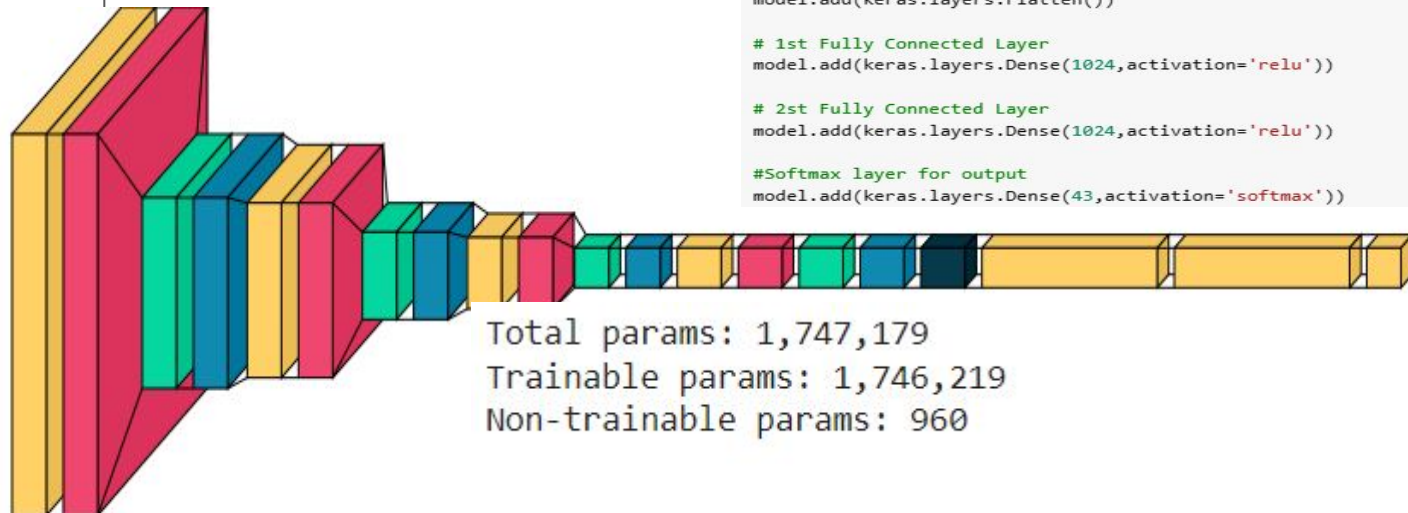
```
model = keras.models.Sequential()
#1st Convolutional Layer
model.add(keras.layers.Conv2D(filters=32, kernel_size=3, activation='relu', input_shape = [50, 50, 3]))
model.add(keras.layers.BatchNormalization())
model.add(keras.layers.MaxPooling2D(pool_size=2))
model.add(keras.layers.Dropout(0.2))
#2nd Convolutional Layer
model.add(keras.layers.Conv2D(filters=64, kernel_size=3, activation='relu'))
model.add(keras.layers.BatchNormalization())
model.add(keras.layers.MaxPooling2D(pool_size=2))
model.add(keras.layers.Dropout(0.2))
#3rd Convolutional Layer
model.add(keras.layers.Conv2D(filters=128, kernel_size=3, activation='relu'))
model.add(keras.layers.BatchNormalization())
model.add(keras.layers.MaxPooling2D(pool_size=2))
model.add(keras.layers.Dropout(0.2))
#4th Convolutional Layer
model.add(keras.layers.Conv2D(filters=256, kernel_size=3, activation='relu'))
model.add(keras.layers.BatchNormalization())
model.add(keras.layers.MaxPooling2D(pool_size=2))
model.add(keras.layers.Dropout(0.2))
#Passing it to a Fully Connected layer
model.add(keras.layers.Flatten())

# 1st Fully Connected Layer
model.add(keras.layers.Dense(1024, activation='relu'))

# 2st Fully Connected Layer
model.add(keras.layers.Dense(1024, activation='relu'))

#Softmax layer for output
model.add(keras.layers.Dense(43, activation='softmax'))
```

Total params: 1,747,179
Trainable params: 1,746,219
Non-trainable params: 960



 Conv2D  BatchNormalization  MaxPooling2D  Dropout  Flatten  Dense


Training :

loss function : “cateorical_crossentropy”. We have trained the model in 15 epochs, but this is a hyper-parameter that we will tune to in the future

Early stopping callback function:

We used the early stopping that is a form of regularization that can be used to prevent overfitting. Is implemented by specifying a callback function in the .fit function


Evaluate the model:



accuracy			0.97	12630
macro avg	0.96	0.95	0.95	12630
weighted avg	0.97	0.97	0.97	12630



Confusion Matrix :

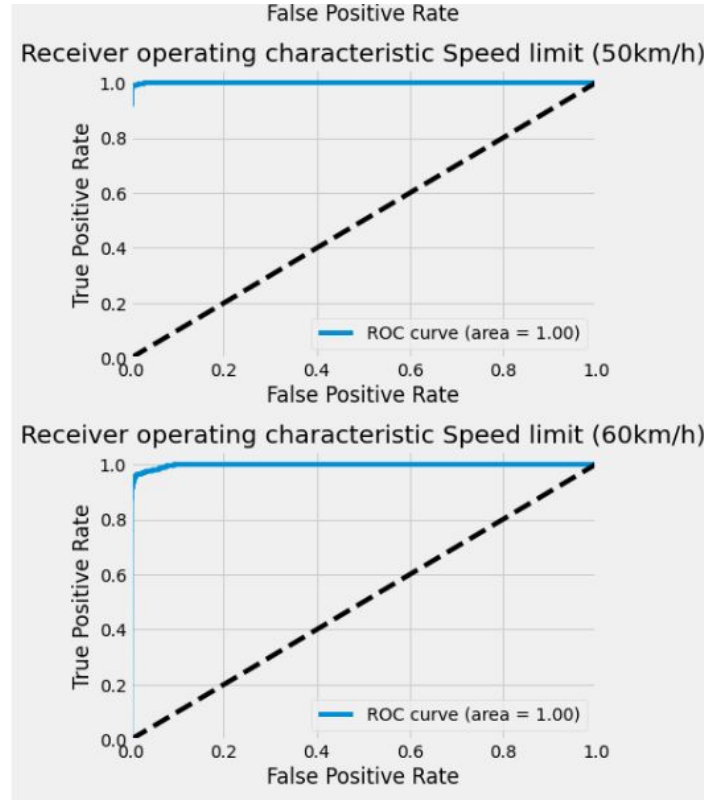


Visual Network Test:

15 random images of the dataset were put to the model test, and we observe how they are classified



ROC Curves (only 2 shown for example):



MNIST Model Based

Architecture:


Stack of alternated Conv2D (with ReLu activation function) and MaxPooling2D layers.

Note that the depth of the feature maps is progressively increasing in the network (from 32 to 256), while the size of the feature maps is decreasing.

We have trained the model in 15 epochs, but this is a hyper-parameter that we will tune to in the future.

Total params: 616,291
Trainable params: 616,291
Non-trainable params: 0

```
model = keras.models.Sequential()  
model.add(keras.layers.Conv2D(filters=32,kernel_size=3,activation='relu', input_shape = [50, 50,3]))  
model.add(keras.layers.MaxPooling2D(pool_size=(2,2)))  
model.add(keras.layers.Conv2D(filters=64 , activation='relu',kernel_size=3))  
model.add(keras.layers.MaxPooling2D(pool_size=(2,2)))  
model.add(keras.layers.Conv2D(filters=128 , activation='relu' ,kernel_size=3))  
model.add(keras.layers.MaxPooling2D(pool_size=(2,2)))  
model.add(keras.layers.Flatten())  
model.add(keras.layers.Dropout(0.5))  
model.add(keras.layers.Dense(250,activation='relu'))  
model.add(keras.layers.Dense(43,activation='softmax'))  
model.summary()
```



 InputLayer  Rescaling  Conv2D  MaxPooling2D  Flatten
 Dense




Training :

loss function : “cateorical_crossentropy”. We have trained the model in 15 epochs, but this is a hyper-parameter that we will tune to in the future

Early stopping callback function:


We used the early stopping that is a form of regularization that can be used to prevent overfitting. Is implemented by specifying a callback function in the .fit function

Evaluate the model:




accuracy			0.98	12630
macro avg	0.96	0.97	0.97	12630
weighted avg	0.98	0.98	0.98	12630



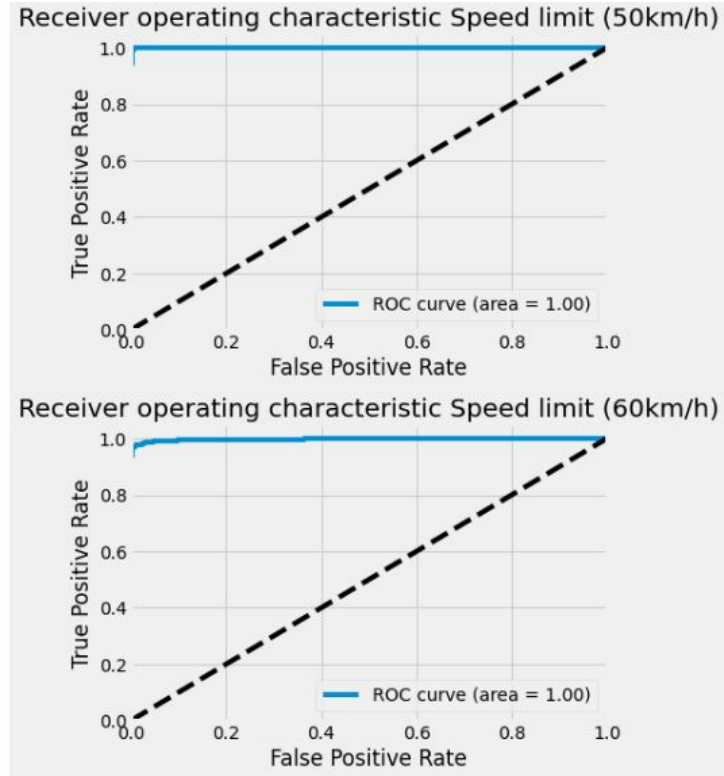


Visual Network Test:

15 random images of the dataset were put to the model test, and we observe how they are classified



ROC Curves (only 2 shown for example):






02

Hyper parameter Tuning



Tuning with keras Tuner



We tried to fine-tune the hyperparameters in order to improve performance using the Keras library called “keras Tuner”.

This process was only carried out for Mnist.

The optimized parameters were:

the number of training epochs, the learning rate, and the number of units in the dense layers of the network.


All other parameters remained unchanged

This are the result of the search:

Best epoch: 37

The optimal number of units in the first densely-connected layer is 288

the optimal learning rate for the optimizer is 0.001






03


CNN USING VCC19



Approach




We used the VGG19 architecture to test a pre-trained network on our dataset through transfer learning. VGG19 is a convolutional neural network (CNN) trained on the ImageNet dataset for image classification. It consists of 19 layers and has been trained to recognize 1000 different classes of objects.



Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 50, 50, 3)]	0
block1_conv1 (Conv2D)	(None, 50, 50, 64)	1792
block1_conv2 (Conv2D)	(None, 50, 50, 64)	36928
block1_pool (MaxPooling2D)	(None, 25, 25, 64)	0
block2_conv1 (Conv2D)	(None, 25, 25, 128)	73856
block2_conv2 (Conv2D)	(None, 25, 25, 128)	147584
block2_pool (MaxPooling2D)	(None, 12, 12, 128)	0
block3_conv1 (Conv2D)	(None, 12, 12, 256)	295168
block3_conv2 (Conv2D)	(None, 12, 12, 256)	590880
block3_conv3 (Conv2D)	(None, 12, 12, 256)	590880
block3_conv4 (Conv2D)	(None, 12, 12, 256)	590880
block3_pool (MaxPooling2D)	(None, 6, 6, 256)	0

```
=====
Total params: 20,024,384
Trainable params: 20,024,384
Non-trainable params: 0
```


Feature Extraction (with data augmentation)



We use feature extraction with a data augmentation layer, freezing all the layers before doing that, and applying the VGG19 convolutional base directly to our model (also applying input scale with `vgg19.preprocess_input`).

```
#freezing all the layers
conv_base.trainable = False
#Setting trainable to False empties the list of trainable weights of the layer or model
print("Trainable weights after freezing", len(conv_base.trainable_weights))

conv_base.summary()
```


Trainable weights after freezing 0
Model: "vgg19"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 50, 50, 3)]	0
block1_conv1 (Conv2D)	(None, 50, 50, 64)	1792
block1_conv2 (Conv2D)	(None, 50, 50, 64)	36928
block1_pool (MaxPooling2D)	(None, 25, 25, 64)	0
block2_conv1 (Conv2D)	(None, 25, 25, 128)	73856
block2_conv2 (Conv2D)	(None, 25, 25, 128)	147584
block2_pool (MaxPooling2D)	(None, 12, 12, 128)	0
block3_conv1 (Conv2D)	(None, 12, 12, 256)	295168
block3_conv2 (Conv2D)	(None, 12, 12, 256)	590080
block3_conv3 (Conv2D)	(None, 12, 12, 256)	590080
block3_conv4 (Conv2D)	(None, 12, 12, 256)	590080

block3_pool (MaxPooling2D)	(None, 6, 6, 256)	0
block4_conv1 (Conv2D)	(None, 6, 6, 512)	1180160
block4_conv2 (Conv2D)	(None, 6, 6, 512)	2359008
block4_conv3 (Conv2D)	(None, 6, 6, 512)	2359008
block4_conv4 (Conv2D)	(None, 6, 6, 512)	2359008
block4_pool (MaxPooling2D)	(None, 3, 3, 512)	0
block5_conv1 (Conv2D)	(None, 3, 3, 512)	2359008
block5_conv2 (Conv2D)	(None, 3, 3, 512)	2359008
block5_conv3 (Conv2D)	(None, 3, 3, 512)	2359008
block5_conv4 (Conv2D)	(None, 3, 3, 512)	2359008
block5_pool (MaxPooling2D)	(None, 1, 1, 512)	0

=====
Total params: 20,024,384
Trainable params: 0
Non-trainable params: 20,024,384
=====

Fine Tuning



For exploiting fine-tuning, we instantiated the vgg19 convolutional base and unfroze the last 5 layers, then applied the densely connected classifier on top of the network
It's not convenient to fine tune all the layers;

```
[ ] # Freezing all layers until the fifth from the last
conv_base.trainable = True
for layer in conv_base.layers[:-5]:
    layer.trainable = False
```

```
[ ] conv_base.summary()
```

Model: "vgg19"



Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 50, 50, 3)]	0
block1_conv1 (Conv2D)	(None, 50, 50, 64)	1792
block1_conv2 (Conv2D)	(None, 50, 50, 64)	36928
block1_pool (MaxPooling2D)	(None, 25, 25, 64)	0
block2_conv1 (Conv2D)	(None, 25, 25, 128)	73856
block2_conv2 (Conv2D)	(None, 25, 25, 128)	147584
block2_pool (MaxPooling2D)	(None, 12, 12, 128)	0
block3_conv1 (Conv2D)	(None, 12, 12, 256)	295168
block3_conv2 (Conv2D)	(None, 12, 12, 256)	590880
block3_conv3 (Conv2D)	(None, 12, 12, 256)	590880

```
block3_conv4 (Conv2D)      (None, 12, 12, 256)      590880
block3_pool (MaxPooling2D) (None, 6, 6, 256)        0
block4_conv1 (Conv2D)      (None, 6, 6, 512)        1180160
block4_conv2 (Conv2D)      (None, 6, 6, 512)        2359808
block4_conv3 (Conv2D)      (None, 6, 6, 512)        2359808
block4_conv4 (Conv2D)      (None, 6, 6, 512)        2359808
block4_pool (MaxPooling2D) (None, 3, 3, 512)        0
block5_conv1 (Conv2D)      (None, 3, 3, 512)        2359808
block5_conv2 (Conv2D)      (None, 3, 3, 512)        2359808
block5_conv3 (Conv2D)      (None, 3, 3, 512)        2359808
block5_conv4 (Conv2D)      (None, 3, 3, 512)        2359808
block5_pool (MaxPooling2D) (None, 1, 1, 512)        0
```


```
=====
Total params: 20,024,384
Trainable params: 9,439,232
Non-trainable params: 10,585,152
```

Highlighted Parameters



- *Loss function: Categorical Cross Entropy*
- *Optimizers: RMSprop vs Adam*
- *Learning Rate Reduction*
- *Batch Normalization layer*
- *Dropout layer*
- *Dense Layer: ReLU activation function*




Performances – Fast Feature Extraction



- RMSprop optimizer



- Adam optimizer




1s 5ms/step - loss: 0.0487 - accuracy: 0.9778 - val_loss: 0.2658 - val_accuracy: 0.9297 - lr: 0.0010


1s 7ms/step - loss: 0.0431 - accuracy: 0.9808 - val_loss: 0.3584 - val_accuracy: 0.9251 - lr: 0.0010



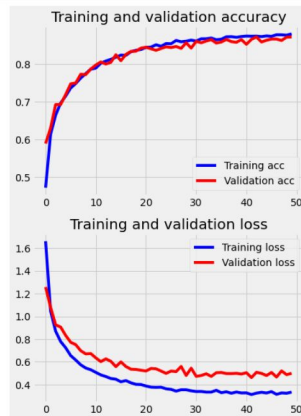
Performances – Feature Extraction with Data Augmentation: training




- RMSprop optimizer




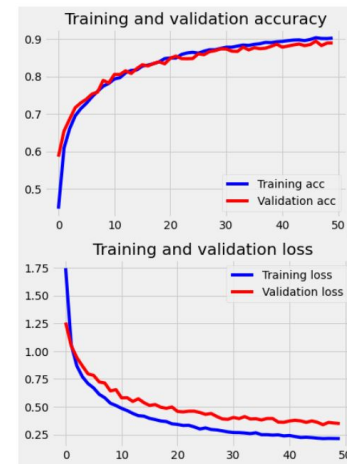
loss: 0.3338 - accuracy: 0.8794 - val_loss: 0.4988 - val_accuracy: 0.8711 - lr: 0.0010




- Adam optimizer



- loss: 0.2131 - accuracy: 0.9023 - val_loss: 0.3480 - val_accuracy: 0.8893 - lr: 0.0010




Performances – Feature Extraction with Data Augmentation: testing





- RMSprop optimizer

395/395 [=====] - 6s 15ms/step
Test Data accuracy: 54.53681710213777




- Adam optimizer


395/395 [=====] - 6s 15ms/step
Test Data accuracy: 53.34125098970704



Performances – Fine Tuning: training





- *RMSprop optimizer*




· loss: 0.1591 - accuracy: 0.9277 - val_loss: 0.1630 - val_accuracy: 0.9452 - lr: 1.0000e-05

- *Adam optimizer*




- loss: 0.0541 - accuracy: 0.9763 - val_loss: 0.1141 - val_accuracy: 0.9629 - lr: 0.0010



Performances – Fine Tuning: testing


- *RMSprop optimizer*

395/395 [=====] - 12s 28ms/step
Test Data accuracy: 76.70625494853523




- *Adam optimizer*

395/395 [=====] - 6s 15ms/step
Test Data accuracy: 82.32779097387174



Comparison between the different approaches and conclusions



- Highest training accuracy with the fast feature extraction approach
- However, it was the most sensitive to overfitting
- With VGG19, the fine tuning with the Adam optimizer gave the best results
- Overall, the best approaches we tried for this task are the ones with training from scratch the AlexNet and the Mnist based architecture models

Possible improvements:

- Try different pre-trained models
- Try a different test set or a different data augmentation layer
- Improve hyperparameters
- Explore other existing techniques to try to carry out this task

